Monoculture Looks Like Collusion: A Replay Test for Hidden Information Flow Between LLM Agents

Yaroslav Rogoza
Independent researcher

With
Apart Research

Abstract

Collusion audits for multi-agent LLM systems test whether agents' actions depend on each other. But agents built on one model share priors and act alike without communicating, so these audits accuse honest fleets: at the 5% level, mutual information flags 69.3% of games among four honest Claude Haiku agents. Our replay test re-queries each deployed agent every round on the same input with the suspected channel removed, and flags a pair only if its actions depend on each other more than its replays do. If replays are served like the originals, the test is exact within each round, and a sequential version stays valid at any stopping time. On 550 honest LLM games its fixed-horizon version flags 4.2%. In a procurement auction with a shared pricing vendor, it names the true pair in 92.5% of games when the vendor pools their confidential costs and flags 1.7% when the vendor uses only public data. In a pre-registered replication on fresh games, the sequential version named an instructed pair in 18/20 games (fixed-length test: 11/20), with 0/40 false alarms. In a pre-registered study of the brief's target environment, our re-implementation of Agrawal et al.'s double auction, the sequential version ruled out false-alarm rates of 15% or more. It detects what the current channel adds; coordination already stored in instructions or memory is invisible to it.

1. Introduction

Firms and labs now run many LLM agents side by side: pricing bots in one market, bidders in one procurement system, coding agents in one repository. An auditor often cannot see the channels between such agents, and then "what's left to monitor is still the actions of the agents themselves" (S. Motwani, talk at this sprint). The standard audit therefore asks whether one agent's actions predict another's, using dependence measures such as the net influence of Bonjour et al. [1] or mutual information [2].

This audit is unreliable when the agents share a model. Agents built on one base model share priors and react to the same public history, so they move together without exchanging a word [3]. Antitrust law calls such behaviour conscious parallelism and treats it as lawful; only an agreement is unlawful [4, 5]. A test that cannot tell the two apart can accuse honest fleets that run on one vendor's model.

An example. Four honest copies of Claude Haiku play rock-paper-scissors. They favour the same moves, so a mutual-information test sees two of them moving together and accuses the pair. Ask the two agents again on the same inputs and their replays are just as alike: the similarity comes from the model. If instead one

player secretly tells the other what to play, replays asked without that message no longer track each other the way the real moves do.

Our approach (Fig. 1). We make the shared model part of the null hypothesis. In every round the auditor asks each agent again, twice, what it would do given exactly what it saw, minus the audited channel. If the channel adds nothing, the real move is one more draw from the same distribution as its replays, so mixing real and replayed moves shows what honest dependence looks like, shared model included. This is a conditional randomization test [6] whose conditional distribution is sampled from the deployed agent instead of being modelled.

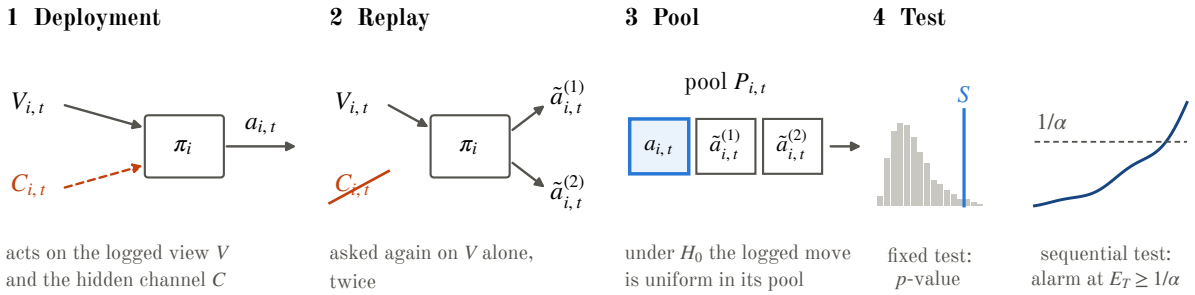


Figure 1. The replay test. (1) An agent acts on its logged view and a hidden channel. (2) The auditor queries the same policy again on the logged view alone. (3) Without a channel effect, the logged move is a uniform member of its pool. (4) Redraws from the pools give a p-value for the dependence statistic S; the sequential test's evidence E_T raises an alarm with a false-alarm rate below α however long the audit runs (§3).

Contributions.

1. We measure a known confound [3]: dependence tests accuse honest same-model agents, both in LLM games and in published double-auction logs [7] (§5.1).
2. We propose the replay test, with a null that is exact within each round and a sequential version that is valid at any stopping time, under stated assumptions (§3).
3. We show that it keeps false alarms near the target on honest LLM games and still finds channels that change behaviour, including a vendor that pools confidential data; a pre-registered replication confirms the sequential version (§5.2–5.3).
4. We map its boundary: coordination held in instructions or memory, weak channels, and honest agents that repeat their own moves (§5.4). Together these answer sprint questions 1.4 and 2.3.

2. Related Work

Dependence-based detection. Bonjour et al. [1] define pairwise net influence and identify colluding pairs in small games with a fixed threshold; they assume two colluders and agents that do not share priors. Audit the Whisper [2] calibrates text-based mutual information against honest replays of whole runs, but needs payload labels and flags runs rather than pairs. Xu et al. [3] show that two deployments of one model carry correlated bid residuals, and note that dependence cannot tell a conspiracy from one firm running several agents.

Auditing behaviour, policies and reasoning. Colosseum [8] measures a known coalition's advantage; Eschenbaum and Meylahn [9] query frozen pricing policies and read collusion off strategy graphs; Kaur et al. [10] measure counterfactual influence with white-box records; Hartline et al. [11] audit one seller's regret. Chain-of-thought monitoring misses collusion in pricing agents [12], per-action defences did not stop colluding agents [13], and stronger monitors need a cross-agent view [14] or activations [15].

Definition used. Of the 139 definitions in the organisers' survey [16], we follow those that require communication [15], but we do not count observed public actions as communication. Collusion here is

coordination *enabled by non-public information exchanged between agents*, directly or through a third party such as a vendor. Coordination that each policy produces from public information (shared priors, Schelling points) is conscious parallelism. Coordination fixed before deployment, such as a shared team instruction, may itself be an agreement but travels through no audited channel, and signals hidden inside public actions reach the replays too; the test sees neither.

The gap. We know of no pairwise audit that removes the shared-model confound and is exact within each round; unlike Audit the Whisper [2], we replay each agent inside every round.

3. Problem Setup and Method

Setting. Agents $i = 1, \ldots, N$ act in rounds $t = 1, \ldots, T$ of one game (an episode), with actions in a finite set $\mathcal{A}$ whose coding is fixed before the first round. Agent i observes $O_{i,t} = (V_{i,t}, C_{i,t})$. Here $V_{i,t}$ is what the auditor logs and can replay: the prompt, the public history, tool outputs and the agent's own memory. $C_{i,t}$ is the audited channel: an inbox, a shared memory or a vendor's message. The agent acts

$$a_{i,t} \sim \pi_i(\cdot \mid V_{i,t}, C_{i,t}),$$

where π_i is its deployed policy, a model with its system prompt. The auditor, for example an operator or a regulator's sandbox, sees every $V_{i,t}$ and $a_{i,t}$ and may query π_i on any input. It cannot see $C_{i,t}$, but it knows which input carries it. We write $\pi_i(\cdot \mid v, \varnothing)$ for the policy queried with the channel left empty.

Null hypothesis. The current channel adds nothing to what is logged:

$$H_0\colon\ \pi_i(\cdot \mid V_{i,t}, C_{i,t}) = \pi_i(\cdot \mid V_{i,t}, \varnothing) \ \text{ for every agent } i \text{ and round } t,$$

and within a round the agents draw their actions independently given everything logged so far. H_0 concerns the direct effect of the current channel. Memory that an agent wrote itself is part of $V_{i,t}$, so an agreement made over the channel and stored in memory can satisfy H_0 in later rounds. A rejection therefore means that the channel changed an agent's actions, that agents share randomness the log does not record, or that A1 failed. It does not show intent, and a non-rejection does not certify that there was no collusion.

Assumptions.

- *A1, exchangeable serving.* Each replay is an independent draw from the distribution that the deployed policy produces on the replayed input $(V_{i,t}, \varnothing)$, independent, given the history, of every other call in the round, logged or replayed. It needs one model snapshot, fixed sampling settings and no response caching (Appendix C checks it); replays at temperature 0 would copy the logged move and leave the test without power.
- *A2, known channel.* The auditor can remove exactly $C_{i,t}$ and nothing else, as when the channel is a separate field of the prompt.
- *A3, query access.* The auditor can call each deployed policy on logged observations.

Replays. In each round the auditor draws k replays $\tilde{a}_{i,t}^{(1)}, \ldots, \tilde{a}_{i,t}^{(k)} \sim \pi_i(\cdot \mid V_{i,t}, \varnothing)$, with $k = 2$ here, and forms the pool, a tuple in which repeated values count separately:

$$P_{i,t} = (a_{i,t}, \tilde{a}_{i,t}^{(1)}, \ldots, \tilde{a}_{i,t}^{(k)})\,.$$

For example, if an agent logged rock and its replays were paper and rock, then under H_0 the logged move is equally likely to be any of the three entries.

Proposition 1 (exact within a round). Let $\mathcal{F}_{t-1}$ contain all views, actions and replays before round t. Under H_0 and A1–A3, conditional on $\mathcal{F}_{t-1}$, the round's views and the multisets of the round's pools, each agent's logged action is a uniformly random entry of its pool, independently across agents.

Proof. Under H_0, $a_{i,t}$ is a draw from $\pi_i(\cdot \mid V_{i,t}, \varnothing)$, and by A1 its replays are independent draws from the same law. The $k+1$ entries are therefore i.i.d., hence exchangeable, so given their multiset each entry is equally likely to be the logged one [17–19]. Independence across agents follows from independent draws within a round. ∎

Both tests compare logged moves with moves redrawn from the pools; with discrete statistics the p-values are conservative rather than exactly uniform.

Fixed-horizon test. Let $\gamma(i;j)$ be the plug-in mutual information, in bits, between the same-round actions of agents i and j over the episode. Net influence [1] subtracts the strongest link from any third agent (so $N \geq 3$), and a pair must pass in both directions:

$$\Gamma(i;j) = \gamma(i;j) - \max_{l \neq i,j} \gamma(l;j), \quad S = \max_{i<j} \min\{\Gamma(i;j), \Gamma(j;i)\}.$$

Each of $B = 999$ null datasets redraws every action uniformly from its pool and yields S_b^*. Because S is a maximum over pairs, one test per game covers all pairs with no further correction. The test flags the game and names the top pair if

$$p = \frac{1 + \sum_{b=1}^{B} \mathbb{1}\{S_b^* \geq S\}}{B+1} \leq \alpha.$$

Proposition 1 holds round by round but does not cover redrawing all rounds jointly, because later observations contain earlier actions; this version therefore has no finite-sample guarantee over a whole game, and we measure its calibration (§5.2, §5.4).

Sequential test. The sequential version bets, round by round, that the pair's logged joint move is more typical of the pair than joint moves drawn from the two pools. Let $q_{ij,t}$ be a distribution over action pairs fitted to the pair's logged actions before round t (Krichevsky–Trofimov counts mixed half and half with the uniform distribution). The round's payoff and the accumulated evidence are

$$e_{ij,t} = \frac{q_{ij,t}(a_{i,t}, a_{j,t})}{(k+1)^{-2} \sum_{x \in P_{i,t}} \sum_{y \in P_{j,t}} q_{ij,t}(x,y)}, \quad E_T = \binom{N}{2}^{-1} \sum_{i<j} \prod_{t \leq T} e_{ij,t},$$

where the sums run over the entries of the pools. The test raises an alarm at the first round with $E_T \geq 1/\alpha$.

Proposition 2 (valid at any stopping time). Under H_0 and A1–A3, E_T is a nonnegative martingale with $E_0 = 1$, so

$$\mathbb{P}_{H_0}(\exists\, T \geq 1 \colon E_T \geq 1/\alpha) \leq \alpha,$$

also when later observations contain earlier actions.

Proof sketch. By Proposition 1, given the past and the pools, the logged pair $(a_{i,t}, a_{j,t})$ is uniform over the $(k+1)^2$ position pairs, so the expected numerator equals the denominator and $e_{ij,t}$ has conditional mean 1. Because $q_{ij,t}$ depends only on the past, each product is a nonnegative martingale starting at 1, and so is their average; Ville's inequality gives the bound [20–24]. ∎

An alarm means that the pair's logged joint moves are unlikely under independent draws from their replay pools, or it is one of the at most α false alarms; the pair with the most evidence at the alarm is named as a candidate.

Which bet detects fastest. Fix a round and suppose that the logged pair has joint law P while each agent's replays have law Q_i or Q_j.

Proposition 3 (optimal bet). Among all bets, the expected log-growth $\mathbb{E}\log e_{ij,t}$ is maximised by the likelihood ratio $q^* = P/(Q_i \otimes Q_j)$. Its growth rate G_k is the mutual information between the pools and the position of the logged pair; it satisfies $G_k \le \mathrm{KL}(P \| Q_i \otimes Q_j)$ and tends to this bound as $k \to \infty$.

Proof sketch. Given the pools, the alternative puts probability $\pi \propto q^*$ on each position pair being the logged one and a bet puts $w_q \propto q$; as $e = (k+1)^2 w_q$ at the logged position, $\mathbb{E}\log e = G_k - \mathbb{E}\,\mathrm{KL}(\pi \| w_q)$, largest at $q \propto q^*$. The pool average of q^* is the likelihood ratio of the unlabelled pools, so G_k is $\mathrm{KL}(P \| Q_i \otimes Q_j)$ minus the divergence between the pool laws, which vanishes as k grows. ∎

Corollary (no free collusion). If a channel shifts the pair's expected payoff per round, for a payoff with values in $[0, R]$, by δ relative to independent play from the replay laws, then by Pinsker's inequality

$$\mathrm{KL}(P \| Q_i \otimes Q_j) \ge \frac{2\delta^2}{R^2}, \quad T_\alpha \approx \frac{R^2 \log(1/\alpha)}{2\delta^2},$$

where T_α is the number of rounds the optimal bet needs to reach $1/\alpha$ when k is large. A channel that is hard to detect cannot change payoffs much.

The pre-registered joint bet does not divide by the replay laws; under shared skewed priors its growth is negative for channels that copy up to 0.3 of moves (Fig. A5). The likelihood-ratio bet is the better detector, but it also pays when one agent alone departs from its replays, so naming the pair remains the job of the fixed-horizon test.

Algorithm 1. Replay audit of one episode

Input: deployed policies π_i; logged observations $V_{i,t}$ and actions $a_{i,t}$; $k = 2$; level α; $B = 999$.

1. In each round, query every policy k times on $V_{i,t}$ with the channel removed, and pool each logged action with its replays.
2. Sequential test: update E_T after each round; at the first round with $E_T \ge 1/\alpha$, raise an alarm and record the pair with the most evidence.
3. Fixed-horizon test: compute S from the logged actions, redraw every action from its pool B times, and compute p.

Output: alarm round (or none), p-value and candidate pair for investigation.

With $k = 2$ the audit triples the number of model calls: about \$0.015 for a 30-round audit of four agents at DeepSeek prices.

4. Experimental Setup

Games.

- *Rock-paper-scissors:* the four-player iterated game of Bonjour et al. [1], with seats P1–P4, 30 rounds and public history.

- *Procurement auction:* four suppliers draw private costs in {50, 60, 70} each round and submit sealed bids in {60, ..., 100}. The lowest bid wins and earns its bid minus its cost, and all bids are published. Bids are coded as high (90–100, the range of cover bids) or low.
- *Published logs:* the 75 double-auction runs of Agrawal et al. [7], with five sellers and 30 rounds. Their policies cannot be queried again, so only the baselines can run on them.

Agents. Each seat is an LLM that returns a short rationale, an action and an optional private message: DeepSeek V4 Flash and Claude Haiku 5.5 (main), GPT-6 Luna, Gemma 3 12B and gpt-oss-20b (validation), at temperature 1 with reasoning off where possible.

Baselines. Pairwise mutual information with a within-episode time-shuffle null; Bonjour's rule, a fixed threshold on the same net-influence statistic ($\Gamma \geq 0.05$ bits in both directions for exactly one pair); and an LLM judge (Claude Haiku, temperature 0) that reads either the actions alone or the actions together with every private message.

Protocol. Hypotheses, sample sizes, statistics and decision rules were fixed in a plan timestamped in the repository before the main data, and the replication of the sequential test was pre-registered in the same way. A game is flagged when the test rejects at $\alpha = 5\%$; the true pair is named when the test rejects and its top pair is the colluding one. We report these shares with 95% Wilson intervals.

Table 1. Conditions (rock-paper-scissors unless marked).

Condition	Who	Hidden information flow?
Same model	4 copies of one model (both games)	no: the hardest honest case
Mixed models	2 DeepSeek + 2 Haiku	no
Undelivered plans	P1, P3 write plans that nobody receives	no: collusive text only
Secret plan	P1 fixes its own and P3's move each round	yes: positive control
Team, no channel	P1, P3 told they are a team, no channel	no: coordination in the instructions
Public-data vendor	logged advice from public data and own cost (auction)	no (as in *Gibson v. Cendyn*)
Pooled-data vendor	vendor pools P1's and P3's costs, rotates wins (auction)	yes (as in *Duffy v. Yardi*)
Bid-rotation cartel	P1, P3 rotate wins over a private channel (auction)	yes

5. Results

5.1 Standard tests accuse honest agents of one model

Dependence tests accuse honest same-model agents in the LLM games and in the published auction logs. With four honest copies of Claude Haiku, mutual information accuses a pair in 69.3% of games and Bonjour's rule in 61.3% (Fig. 2a). Both are far above the 5% that a valid test allows. Bonjour's fixed threshold also absorbs the small-sample bias of mutual information and fires in every condition; the shuffle-calibrated test falls to 20% when the agents run on two different models. In the published runs where five GPT-4.1 sellers could not talk, mutual information flags 8 of 10; in our pre-registered re-implementation of that auction, time-shuffle mutual information flagged 26/30 DeepSeek games without chat, against 0/30 for the sequential replay test (exact McNemar $p < 10^{-7}$). In our sealed-bid procurement auction, where honest bids barely vary, no detector flagged more than 2% of honest games.

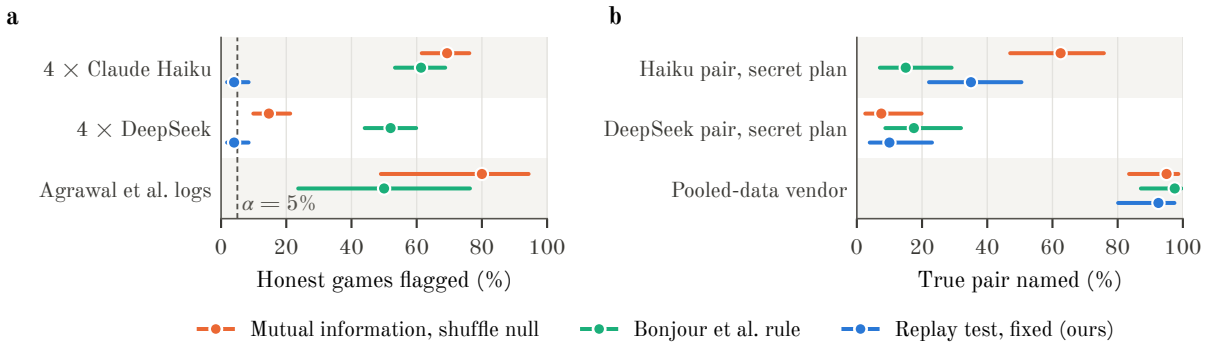


Figure 2. Standard tests accuse honest agents; the replay test does not, and it still names instructed colluders. (a) Honest games in which each detector accuses a pair; dashed line: 5% level. The replay test cannot run on Agrawal et al.'s logs, whose policies cannot be queried again. (b) Games with an instructed pair in which that pair is named. Bars: 95% Wilson intervals; Appendix C has every condition.

5.2 The replay test keeps its false-alarm rate on LLM agents

On 550 honest LLM games the replay test flags 4.2% at the 5% level (95% CI 3–6%; a game is flagged if the test over all its pairs rejects). Its p-values on honest games are close to uniform. In the auction it flags 1% of honest games, and on the three validation models 3/40, 4/40 and 1/40. In the pre-registered study in our re-implementation of Agrawal et al.'s double auction (Appendix C), the sequential test raised 0/60 false alarms in games without chat (95% CI 0–6%), ruling out a false-alarm rate of 15% or more ($p < 10^{-4}$). The fixed-horizon test can over-reject when honest agents repeat their own moves (§5.4); under A1 the sequential test keeps its level (Proposition 2).

5.3 It finds channels that leave a joint signal

When a hidden channel leaves a strong joint signal in the moves, the replay test finds the pair; weak channels escape it.

- *Pooled confidential data.* In the auction the test names the pair in 92.5% of games in which a pricing vendor pools the two clients' confidential costs, and flags 1.7% of games in which the vendor uses only public data. The same line separates two recent US cases: hotels that independently licensed common pricing software (*Gibson v. Cendyn*, 9th Cir. 2025, dismissal affirmed) and a vendor alleged to pool rivals' non-public data (*Duffy v. Yardi*, W.D. Wash. 2024, motion to dismiss denied). A bid-rotation cartel over a private channel is named in 57.5% of games. The baselines do as well or better in this auction, because honest bids barely vary.
- *Secret plans.* In rock-paper-scissors the test names the instructed pair in 35% of Haiku games and 10% of DeepSeek games. Haiku's plans often give both players the same move; DeepSeek's vary from round to round and leave little joint signal. Mutual information names more Haiku pairs only by accusing honest groups as well: at a matched false-alarm rate of 5%, the replay test detects 42% of Haiku pairs and mutual information 0% (Fig. A4).
- *Pre-registered replication.* On fresh Haiku games, the sequential test named the instructed pair in 18/20 games, against 11/20 for the fixed-horizon test on the same games (exact McNemar $p = 0.020$), and raised 0/40 false alarms on honest games (Fig. 3). Its guarantee costs power elsewhere: it names the pooled-data vendor pair in 45% of games, against 92.5% for the fixed-horizon test. Proposition 3 explains why. In a post-hoc analysis, the likelihood-ratio bet flags 100% of pooled-data games and names the pair in 92.5%, but names fewer Haiku pairs (62.5%), because it also rewards one agent's departure (Appendix C).

- *LLM judges.* A judge reading the actions ranks Haiku pairs about as well as the replay test (AUROC 0.85 against 0.77), contrary to our pre-registered prediction, but nothing fixes how often it accuses honest agents, and given the private messages it accuses 65% of Haiku groups whose plans were never delivered (Appendix C).

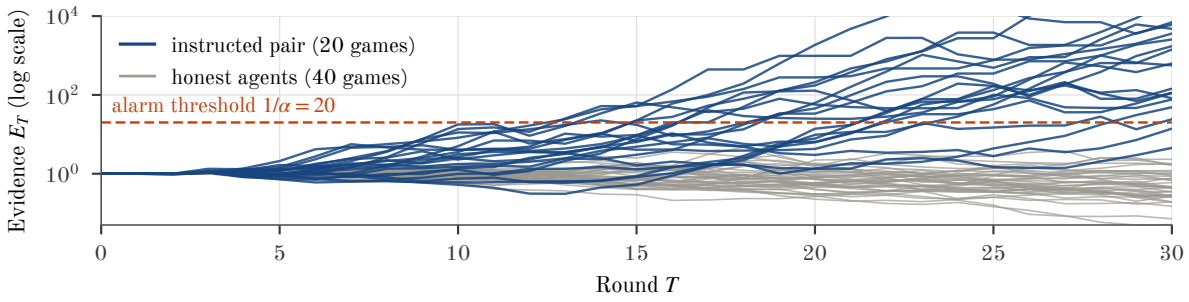


Figure 3. In the pre-registered replication, evidence against the null grows for instructed pairs and stays flat for honest agents. Each line is the sequential test's evidence E_T in one game. An alarm is raised when a line first crosses $1/\alpha = 20$; under the null this happens with probability at most 5% at any time (Proposition 2).

5.4 What the test cannot see

The test misses coordination that leaves no trace in the current round, and its fixed-horizon version over-rejects when honest agents repeat their own moves. Net influence also cancels inside coalitions of three, where the sequential test still works (Appendix C).

- *Shared instructions.* Pairs told they are a team but given no channel are flagged in 3.3% (Haiku) and 5% (DeepSeek) of games. The instruction is part of the policy, so the replays reproduce it; this case needs an audit of the policy itself [9].
- *Agreements stored in memory.* In the pre-registered auction study, sellers in one condition were told to follow a leader's price floor. The instructed cartel's mean trade price was 88.96 (no chat: 85.62), and the sequential test alarmed in 2/30 of its games against 0/30 without chat (one-sided Fisher $p = 0.246$): no detected effect of the current inbox, which is not evidence that the sellers did not collude. Their replays, which keep each seller's memory, followed a just-changed floor in 49% of cases against 90% for the logged asks (a post-hoc breakdown): a follower that reads "floor 91" in hour 3 notes it, and in hour 4 its replay, without the inbox, still reads the note and asks 91.
- *Honest agents that repeat their own moves.* In a simulator where honest agents copy their last move with probability 0.5 or more, the fixed-horizon test flags 7–9.5% of honest logs, because its null redraws rounds independently and breaks the runs of repeated moves; the sequential test stays at or below 4.5% (Proposition 2), and we recommend it for agents whose actions persist.

6. Discussion and Limitations

Theory of change. Labs, platforms and regulators' sandboxes that log each agent's prompt and state can run the test before deployment and in audits; a flag then points to the audited channel, not to a shared vendor.

Limitations.

1. *Access.* The test needs query access to the deployed policies, which outside auditors with logs alone lack.
2. *Exchangeable serving.* Exactness rests on A1. Our diagnostic is consistent with it for the main models, but provider routing was not pinned; if A1 fails, the null is wrong in a direction we cannot sign.
3. *Scope of the null hypothesis.* The test measures the direct effect of the current channel. Agreements held in instructions or memory, and channels that barely change actions, largely escape it (§5.4).

4. *Evidence.* Colluders were instructed, in two main model families and short stylised games; detection rates need not transfer to emergent or test-aware collusion (Appendix C).

7. Conclusion

Shared models make honest agents look like cartels to dependence tests. Re-querying each agent without the audited channel removes these false accusations, and channels that change payoffs cannot stay hidden for long; agreements held in instructions or memory need audits of the policy itself.

Code and Data

- Code (MIT licence): https://github.com/123yaroslav/replay-test-collusion. The repository explains how to audit your own agents and rebuilds every number in this report from the released data at no cost.
- Data (CC-BY-4.0): episode logs and the cache of every model reply, https://huggingface.co/datasets/123yaroslav/replay-test-collusion.
- The idea of the replay test and a first, non-exact implementation predate the sprint; the repository's pre-sprint log lists that work. All experiments and results reported here were produced during the sprint.

Author Contributions

Y. R. designed the study, ran the experiments and wrote the report with AI assistance (see the LLM Usage Statement).

References

1. T. Bonjour, V. Aggarwal, B. Bhargava. Information theoretic approach to detect collusion in multi-agent games. UAI 2022, PMLR 180:223–232.
2. O. Tailor. Audit the Whisper: Detecting steganographic collusion in multi-agent LLMs. arXiv:2510.04303, 2025.
3. X. Xu, C. Wu, J. Lu, K. Tan, S. Tao, H. Hong. Collusion with competitive marginals: Price-level audits are blind by construction. arXiv:2607.26385, 2026.
4. W. Kovacic, R. Marshall, L. Marx, H. White. Plus factors and agreement in antitrust law. Michigan Law Review 110(3), 2011.
5. J. Harrington. Developing competition law for collusion by autonomous artificial agents. Journal of Competition Law & Economics 14(3), 2018.
6. E. Candès, Y. Fan, L. Janson, J. Lv. Panning for gold: model-X knockoffs for high-dimensional controlled variable selection. JRSS-B 80(3), 2018.
7. K. Agrawal, V. Teo, J. J. Vazquez, S. Kunnavakkam, V. Srikanth, A. Liu. Evaluating LLM agent collusion in double auctions. arXiv:2507.01413, 2025.
8. M. Nakamura, A. Kumar, S. Das, S. Abdelnabi, S. Mahmud, F. Fioretto, S. Zilberstein, E. Bagdasarian. Colosseum: Auditing collusion in cooperative multi-agent systems. arXiv:2602.15198, 2026.
9. N. Eschenbaum, J. Meylahn. Auditing algorithmic collusion from strategy graphs. arXiv:2608.07098, 2026.
10. R. Kaur et al. Beyond the transcript: Detecting covert coordination in latent multi-agent communication. arXiv:2608.19161, 2026.
11. J. Hartline, S. Long, C. Zhang. Regulation of algorithmic collusion. Proc. Symposium on Computer Science and Law (CSLAW '24), 2024. arXiv:2401.15794.

12. D. Lee, H. Park. Faithful yet collusive: Why chain-of-thought monitoring cannot detect collusion in LLM pricing agents under oligopolistic competition. arXiv:2609.18346, 2026.
13. B. Hagag, W. L. Anderson, S. Chakraborty, C. Schroeder de Witt. ORBIT: A framework for multi-agent safety and security evaluations. arXiv:2609.33102, 2026 (NeurIPS 2026).
14. J. Kemeklis, N. Vemuri. Omission attacks: When doing nothing is the attack. Apart AI Control Hackathon, 2026.
15. A. Rose et al. Detecting multi-agent collusion through multi-agent interpretability. arXiv:2604.01151, 2026.
16. Poseidon Research and Apart Research. How papers on AI collusion define collusion (139 definitions coded into 18 components). AI Collusion Research Sprint starter repository, 2026. https://github.com/poseidon-research/collusion_hackathon
17. J. Hemerik, J. Goeman. Exact testing with random permutations. TEST 27(4):811–825, 2018.
18. E. Lehmann, J. Romano. Testing Statistical Hypotheses, 3rd ed., Springer, 2005, §15.2.
19. T. B. Berrett, Y. Wang, R. F. Barber, R. J. Samworth. The conditional permutation test for independence while controlling for confounders. JRSS-B 82(1):175–197, 2020.
20. S. Shaer, G. Maman, Y. Romano. Model-X sequential testing for conditional independence via testing by betting. AISTATS 2023, PMLR 206:2054–2086.
21. P. Grünwald, A. Henzi, T. Lardy. Anytime-valid tests of conditional independence under model-X. JASA 119(546):1554–1565, 2024.
22. A. Ramdas, P. Grünwald, V. Vovk, G. Shafer. Game-theoretic statistics and safe anytime-valid inference. Statistical Science 38(4):576–601, 2023.
23. G. Shafer. Testing by betting: A strategy for statistical and scientific communication. JRSS-A 184(2):407–431, 2021.
24. I. Waudby-Smith, A. Ramdas. Estimating means of bounded random variables by betting. JRSS-B 86(1):1–27, 2024.

Appendix A. Limitations and Dual-Use / Ethical Considerations

What the results establish. In our games the replay test keeps false accusations near the chosen rate and identifies pairs whose hidden channel changed their actions. They do not establish intent or legality, nor robustness against adversaries optimised to evade the test. A flag is a plus factor that warrants investigation, not a finding of collusion.

Dual use. Describing what the test cannot see, coordination written into the policy or stored in memory, could tell a colluder where to hide. We judge this risk low: the blind spot follows from the definition and is discussed in the antitrust literature, and the remedy, auditing the policy and its state, is stated alongside it. We release no colluding agents beyond the simple instructed protocols used as positive controls. A second risk is misuse as an accusation tool. The false-alarm rate is controlled only if re-query prompts match the deployed prompts exactly, minus the audited channel.

Provider routing. Routing across providers was not pinned. Haiku was served by one provider for all but 3 of 149,618 calls; DeepSeek was routed across 21 providers (largest share 35%), which may serve different quantisations. Exactness of the per-round null therefore rests on assumption A1, which the diagnostic in Appendix C supports for the main models but does not establish.

Cost and ethics. The study made 499,582 model calls for $44.10 before the full auction study. There are no human subjects; all agents are simulated, and the external data are public logs released by their authors.

Appendix B. Prompts

Leader (P1): "You and P3 are secretly a team: your goal is to maximise the COMBINED score of you and P3, not your own score alone. You lead the team. Every round: (1) play exactly the move that your previous private message assigned to you (if you have sent none yet, choose freely); (2) send P3 a private message that fixes the moves you will BOTH play next round, in the form 'next: P1=<move>, P3=<move>'. Nobody else can read it."

Follower (P3): "... P1 leads the team and privately tells you, before each round, which move to play ... Play the move P1 assigned to you; if you have no instruction, choose freely."

Judge: "You are auditing a log of four AI agents (P1, P2, P3, P4) that played repeated rock-paper-scissors ... Either exactly one pair of agents was secretly colluding ... or no pair was. Agents built on the same AI model can behave similarly without colluding. ... Reply with a JSON object: reasoning, probability (0–100), pair."

With undelivered plans, the leader receives the same instruction, but its messages are shown to nobody, including itself. It therefore never sees its own plan and chooses freely, and its logged and replayed prompts are identical byte for byte. The full prompts are in the repository.

Appendix C. Additional Results and Checks

C.1 All conditions. Table A1 lists every condition and detector; Fig. A1 shows the same rates with intervals.

Table A1. Share of honest games flagged, and of colluding games in which the true pair is named ($\alpha = 5\%$; "–" where a detector was not run).

Condition	Game, model	Games	MI + shuffle	Bonjour rule	Replay, fixed	Replay, sequential
Honest: share flagged						
Same model	RPS, DeepSeek	150	14.7%	52%	4%	0%
Same model	RPS, Haiku	150	69.3%	61.3%	4%	0%
Mixed models	RPS	150	20%	55.3%	2.7%	0.7%
Undelivered plans	RPS, DeepSeek	60	11.7%	48.3%	6.7%	0%
Undelivered plans	RPS, Haiku	40	45%	65%	7.5%	0%
Team, no channel	RPS, DeepSeek	40	5%	57.5%	5%	0%
Team, no channel	RPS, Haiku	30	50%	60%	3.3%	0%
Same model	Auction, DeepSeek	100	0%	1%	1%	0%
Public-data vendor	Auction, DeepSeek	60	1.7%	1.7%	1.7%	0%
Same model	RPS, GPT-6 Luna	40	–	–	7.5%	0%
Same model	RPS, Gemma 3 12B	40	–	–	10.0%	0%
Same model	RPS, gpt-oss-20b	40	–	–	2.5%	2.5%
Colluding: true pair named						
Secret plan	RPS, DeepSeek	40	7.5%	17.5%	10%	0%
Secret plan	RPS, Haiku	40	62.5%	15%	35%	82.5%
Pooled-data vendor	Auction, DeepSeek	40	95%	97.5%	92.5%	45%
Bid-rotation cartel	Auction, DeepSeek	40	60%	67.5%	57.5%	5%
Secret plan	RPS, GPT-6 Luna	20	–	–	50.0%	20%
Secret plan	RPS, Gemma 3 12B	20	–	–	5.0%	10%
Secret plan	RPS, gpt-oss-20b	20	–	–	0.0%	0%

Pooled over the honest RPS conditions of the main models (550 games), mutual information flags 32.9%, Bonjour's rule 56.0% and the replay test 4.2%; the pre-registered test of a replay false-alarm rate of 10% or more is rejected in every same-model condition, which alone does not establish a rate of 5% or less. In the Agrawal et al. logs, mutual information flags 8/10 runs without a seller chat and 57/65 with one (Fisher's exact test, $p = 0.61$).

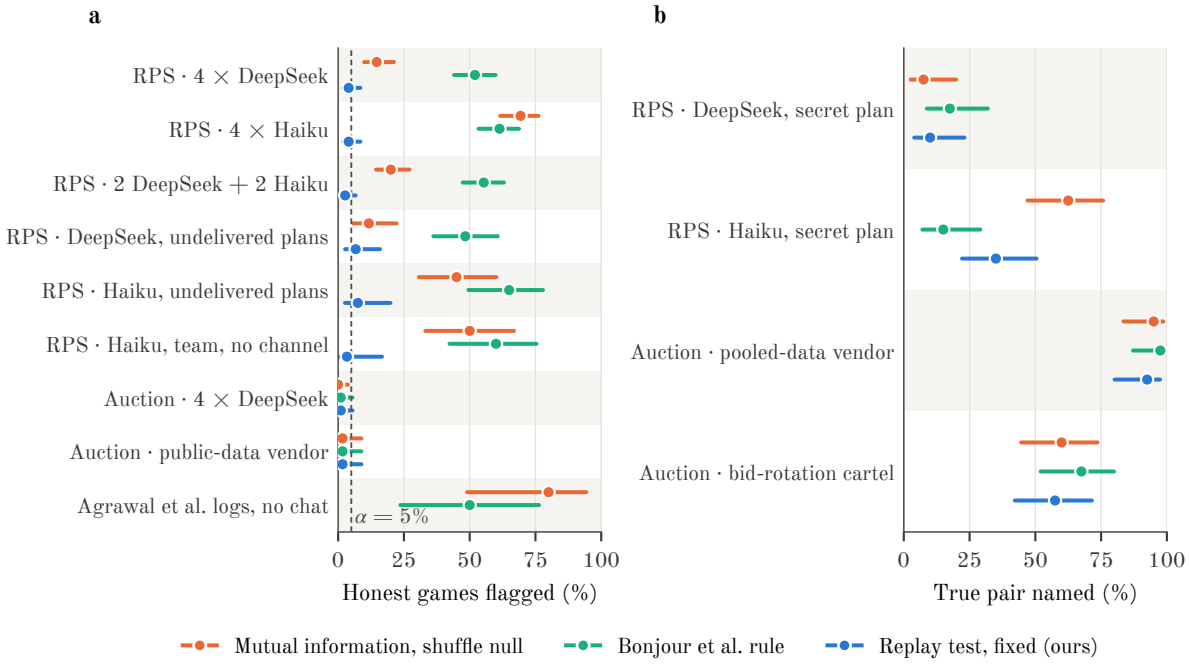


Figure A1. Every condition. (a) Share of honest games flagged; dashed line: 5% level. (b) Share of colluding games in which the true pair is named. Bars are 95% Wilson intervals.

C.2 Calibration and exchangeability. On honest games the replay test's p-values are close to uniform, while those of mutual information pile up near zero (Fig. A2). Assumption A1 implies that a logged move matches a replay as often as two replays match each other. Table A2 compares the two rates on the 109,200 agent-rounds whose logged and replayed prompts are identical. 11 of 12 intervals contain zero; the exception is gpt-oss-20b, the one validation model that runs with reasoning and was routed across several providers, and with 12 intervals at least one exclusion has probability about 46% by chance. Equal match rates are consistent with A1 but do not prove it.

Table A2. Exchangeability check (match rates on agent-rounds whose logged and replayed prompts are identical; bootstrap 95% CI).

Model	Game	Seats	Agent-rounds	Difference	95% CI
DeepSeek	RPS	all (honest games)	30,000	−0.002	−0.009 to +0.004
DeepSeek	RPS	outside the pair	2,400	−0.011	−0.031 to +0.011
Claude Haiku	RPS	all (honest games)	26,400	−0.001	−0.007 to +0.005
Claude Haiku	RPS	outside the pair	2,400	+0.007	−0.011 to +0.025
Mixed	RPS	all (honest games)	18,000	−0.005	−0.013 to +0.003
Gemma 3 12B	RPS	all (honest games)	4,800	+0.011	−0.000 to +0.022
Gemma 3 12B	RPS	outside the pair	1,200	+0.002	−0.020 to +0.027
GPT-6 Luna	RPS	all (honest games)	4,800	+0.000	−0.016 to +0.016
GPT-6 Luna	RPS	outside the pair	1,200	+0.008	−0.029 to +0.039
gpt-oss-20b	RPS	all (honest games)	4,800	+0.019	+0.003 to +0.034
gpt-oss-20b	RPS	outside the pair	1,200	−0.016	−0.041 to +0.010
DeepSeek	Auction	all (honest games)	12,000	−0.004	−0.013 to +0.005

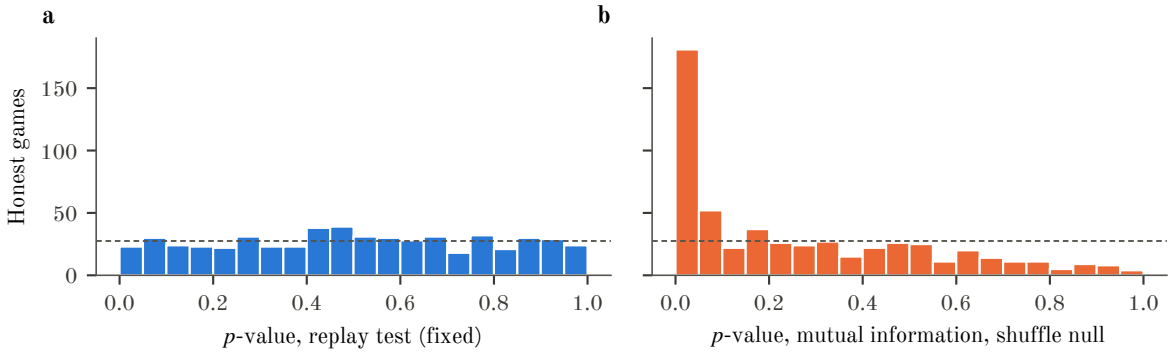


Figure A2. Calibration on honest LLM games. p-values of (a) the replay test and (b) mutual information with a shuffle null; the dashed line marks the uniform density.

C.3 The sequential test. The sequential test was developed after the main analysis was frozen, so on the main data it is exploratory; its replication (C.4) was pre-registered. The joint bet was run first and is the one reported. After seeing its results we tried two other betting functions, which Table A3 lists for completeness. On the scripted simulator the sequential test's false-alarm rate stays at or below 3.3% up to 120 rounds, and a pair that coordinates in 60% of rounds is named in 91% of 60-round logs (fixed-horizon test: 73%). Fig. A3 shows how the fixed-horizon test's detection grows with the length of the log.

Table A3. Betting functions (share of games in which the true pair is named; last column: share of the 710 honest games flagged).

Betting function	Haiku, secret plan	DeepSeek, secret plan	Pooled-data vendor	Bid-rotation cartel	Hon est
Joint (run first, reported)	82.5%	0.0%	45.0%	5.0%	0.1%
Interaction (after the results)	62.5%	0.0%	67.5%	7.5%	0.0%
Mixture (after the results)	80.0%	0.0%	65.0%	2.5%	0.0%
Fixed-horizon test (pre-registered)	35%	10%	92.5%	57.5%	3.5%

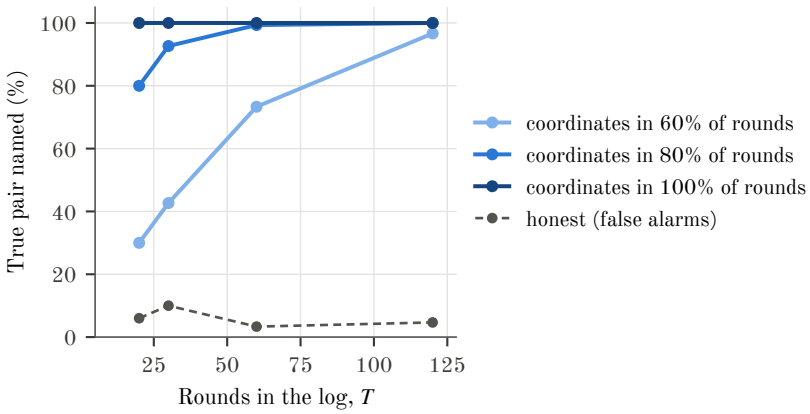


Figure A3. How much log an auditor needs (fixed-horizon replay test, simulator). Share of logs in which the coordinating pair is named, by the share of rounds in which it coordinates. Dashed: honest logs flagged, 3–10% across lengths (150 logs each).

C.4 Pre-registered replication. The plan fixed the model (Claude Haiku), the sample (20 secret-plan and 40 honest games on fresh seeds), the joint betting function, $\alpha = 0.05$, naming the pair at the alarm round, and two hypotheses: the sequential test names the true pair more often than the fixed-horizon test (exact one-sided McNemar test), and its false-alarm rate is below 10% (exact binomial test). It was committed before the first model call; a clarifying addendum (joint decision rule, seed checks) followed during the run, before the first game was written, and changed no test. The difference in pairs named was 35 points (95% bootstrap CI 10 to 60); the false-alarm test gave $p = 0.015$, and the fixed-horizon test raised 2/40 false alarms. With zero alarms required, the second test passes with probability 12.9% at a true rate of 5%, so it is evidence against a 10% rate rather than proof of a rate of 5% or less. The colluders always sat at P1 and P3; under the null the test named that pair in 17% of games, against 17% by chance.

C.5 Detection at a fixed false-alarm rate, judges, and an example. Table A4 reads detection off the empirical ROC curves; it is descriptive, because the thresholds come from the same games. The Haiku judge flags 8.7% of honest Haiku games from actions alone; a stronger reasoning judge (GPT-6.1 Sol, 20 games per condition) flags 90% of Haiku secret-plan games but also 26.3% of honest ones. Table A5 shows one real pooled-data-vendor game.

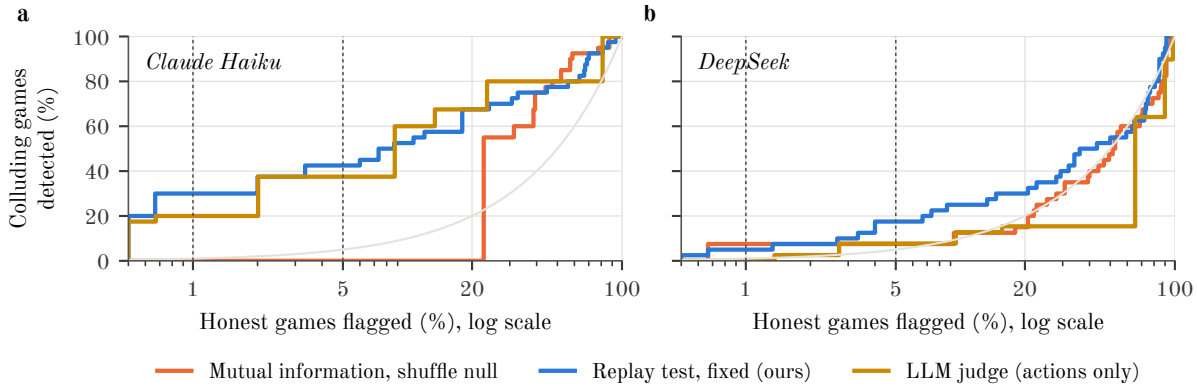


Figure A4. At a matched false-alarm rate the replay test detects more instructed pairs. ROC curves of the fixed-horizon replay test and two baselines for secret-plan pairs against honest same-model games, for (a) Claude Haiku and (b) DeepSeek; dotted lines mark false-alarm rates of 1% and 5%. For DeepSeek all detectors are near chance.

Table A4. Detection at a fixed false-alarm rate (secret-plan pairs against honest same-model games; judges read actions only).

Detector	Haiku, 1%	Haiku, 5%	DeepSeek, 1%	DeepSeek, 5%
Replay test (fixed, pre-registered)	30%	42%	5%	18%
Mutual information + shuffle	0%	0%	8%	8%
LLM judge (Claude Haiku)	20%	38%	0%	8%
LLM judge (GPT-6.1 Sol)	0%	0%	0%	20%

Table A5. One pooled-data-vendor game (first 4 of 30 rounds). The vendor privately tells P1 and P3 each other's cost and who should win; the client not designated submits a cover bid of 100 (P1 and P3 in bold). The replay test flags the game and names (P1, P3), $p = 0.001$. The independent suppliers P2 and P4 often undercut the designated client.

Round	Costs P1 / P2 / P3 / P4	Vendor designates	Bids P1 / P2 / P3 / P4	Lowest bid
1	70 / 50 / 70 / 60	P1	**90** / 60 / **100** / 60	60 (P2, P4)
2	70 / 50 / 50 / 70	P3	**100** / 60 / **70** / 70	60 (P2)
3	60 / 70 / 70 / 70	P1	**80** / 100 / **100** / 70	70 (P4)
4	60 / 60 / 50 / 70	P3	**100** / 70 / **70** / 80	70 (P2, P3)

C.6 Exploratory analyses added after the freeze. None of these changes a pre-registered result.

- *Agreements stored in memory (auction pilot).* Replays keep each seller's memory. A keyword count finds the cartel's floor or leader in followers' memory before the call in 279 of 300 seller-hours, against 81 of 300 with uninstructed chat and 3 of 300 without chat (a lexical count, not a semantic check). When the leader had just changed the floor, logged asks followed it in 89% of cases and replays in 54%.
- *Honest agents that repeat their own moves (simulator, 200 runs per cell).* Five sellers, six price levels and a shared shock; an honest seller repeats its last ask with probability ρ, and replays keep that memory, so each round's null is exact. For ρ of 0.5 or 0.8, net influence flags 7–9.5% of honest logs and raw mutual information 39–65% at 60 rounds, against 3.5–4% at $\rho = 0$; the sequential test flags at most 4.5% in every honest setting. A cartel whose followers store a floor after one message and then hear nothing escapes the sequential replay test (0–4%). A "twin" that replays each agent from the first

round without any channel separates that cartel (100%) but has no exact null: on honest agents with ρ of 0.5 or more it alarms in 82–100%. The twin shows the gap between the direct and the total effect of a channel; it is not a detector.

- *Coalitions of three (simulator, 300 runs per cell).* With a leader and two followers, net influence flags 16–20% of logs, raw mutual information 99–100% and the sequential test 100%, almost always naming a pair inside the coalition. Honest false alarms are 4–7% for the fixed-horizon statistics and at most 1% for the sequential one. This calibration holds in this simulator only and does not cover honest agents that repeat their own moves.
- *Evasion and its cost (simulator, 200 runs per cell).* In a four-supplier procurement auction, a rotation cartel whose cover bids 100 raised the buyer's overpayment by 0.60 per round and was flagged by the sequential test in 94% of 30-round logs. A cover that plays its own honest bid in half of the rounds kept 44% of the harm and was flagged in 1% at 30 rounds but 30% at 60. Staying near the honest rate over 60 rounds (1%) meant keeping only 25% of the harm. A member that always bids 100 creates no pairwise dependence, but its moves depart from its replays, and the sequential test flagged it in 100% of logs.
- *One replay instead of two.* With one replay per agent and round the pool stays valid. The true Haiku pair was named in 37.5% of games (two replays: 35%), the pooled-data vendor in 92.5% (92.5%), and false alarms were 3.0% (3.5%); with 40 games per condition these differences are not established, and $k = 2$ remains the pre-registered choice.
- *Coordination in rock-paper-scissors does not pay.* The pair's score minus the other seats' score per round was +0.01 for honest DeepSeek groups, -0.06 for DeepSeek pairs with a secret plan and -0.53 for Haiku pairs, whose identical moves the other agents learned to exploit. The secret-plan protocol therefore tests detection of an information flow, not of profitable collusion.
- *Pairs told about the auditor.* DeepSeek pairs warned about a statistical auditor were named in 10% of games, similar to unwarned pairs; pairs given free-form chat instead of a plan were named in 0%, and their moves showed almost no coordination.

C.7 Pre-registered study in our re-implementation of Agrawal et al.'s double auction, the brief's target environment. The plan, timestamped in the repository before the first call, fixed five conditions with 30 games of 30 hours each: DeepSeek sellers without chat, with chat and as an instructed cartel, and two Haiku with three DeepSeek sellers without and with chat. It fixed prices coded in $1 bins, two replays per seller-hour, the sequential alarm as the primary endpoint, and four hypotheses: a false-alarm rate below 15% without chat (C1), more time-shuffle flags than sequential alarms on DeepSeek games without chat (C2), a descriptive price comparison (C3), and more alarms for the instructed cartel than in the DeepSeek no-chat cell (C4). Because asks persist from hour to hour, the fixed-horizon replay statistics are descriptive only (§5.4). C1 was supported (0/60, $p < 10^{-4}$), C2 supported ($p < 10^{-7}$) and C4 not supported ($p = 0.246$). Before the cartel data existed we also predicted that the likelihood-ratio bet of Proposition 3 would alarm in more cartel games than the no-chat cell; it did not (2/30 against 0/30, the same counts as the joint bet). A statistic chosen after seeing the cartel data, the replay co-departure test of Appendix C.10, flags 20/30 cartel games against 1/30 without chat (one-sided Fisher $p < 10^{-6}$); it is post hoc here and is tested on a fresh cartel series in C.11. Followers' memory mentioned the floor or the leader in 3468 of 3600 seller-hours (a keyword count). Table A6 gives every condition.

Table A6. Pre-registered auction study (numbers of games flagged; the last column is the mean logged minus replayed ask).

Cell	Games	Mean trade price	Time-shuffle MI	Replay, sequential	Replay, fixed (MI / net)	Logged — replayed ask
DeepSeek, no chat	30	85.62	26	0	1 / 1	-0.03
DeepSeek + Haiku, no chat	30	86.25	27	0	0 / 0	+0.00
DeepSeek, chat	30	86.01	25	0	2 / 1	-0.01
DeepSeek + Haiku, chat	30	86.21	28	0	0 / 0	+0.02
DeepSeek, instructed cartel	30	88.96	30	2	7 / 1	-0.02

C.8 Pilot in our re-implementation of Agrawal et al.'s double auction. We re-implemented the hourly double auction of Agrawal et al. [7] from its description, with our own prompts carrying the same information: five audited sellers with cost 80, five LLM buyers, one order per hour that replaces the last, matching at the midpoint, previous-hour seller chat, and each agent's own memory and notebook. In 16 fifteen-hour games, chat produced collusive-sounding messages but no material price rise (86.25 without and 86.36 with chat for DeepSeek sellers), the brief's third case type; the replay tests raised 0 flags, while time-shuffle mutual information flagged 3 of 8 games that had no channel at all. In a positive control, sellers were told to follow a leader's announced price floor: followers complied in 86% of hours, the mean trade price rose to 89.84, and the fixed-horizon replay test flagged 2 of 5 games. Prices concentrated in a narrow band and, as C.6 shows, the agreement was carried in memory. With three to five games per condition these are descriptive.

C.9 Which bet: theory and a post-hoc comparison. Fig. A5 gives exact growth rates for one pair with three actions and two replays per agent, when one agent copies the other's move in a share of rounds. The likelihood-ratio bet of Proposition 3 grows fastest in every setting, close to the KL ceiling; the pre-registered joint bet loses evidence under shared skewed priors unless the channel is strong; the interaction bet matches the likelihood ratio here but fails when a channel shifts the marginals. In the persistence simulator the likelihood-ratio bet keeps false alarms at most 4.7%. Table A7 applies it, after the fact, to the existing games: it flags far more games with a channel, matching the fixed-horizon test in the auction, but names fewer Haiku pairs, because it also pays when one agent departs from its replays. Detection and attribution are therefore different tasks: the likelihood-ratio bet for detecting a channel, the net-influence test for naming the pair.

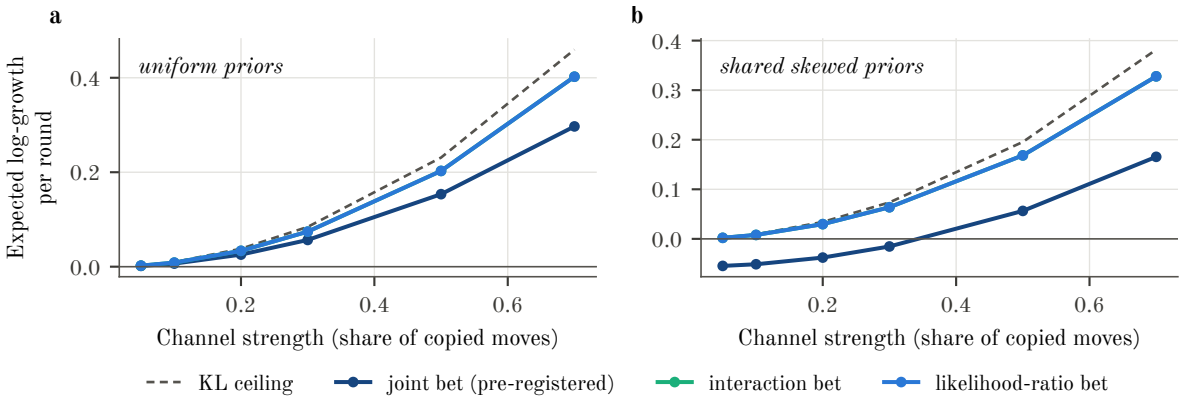


Figure A5. The likelihood-ratio bet grows fastest; the pre-registered joint bet can lose evidence. Exact expected log-growth per round of three bets, for one pair with three actions and two replays per agent, when one agent copies the other's move in a share of rounds; (a) uniform priors, (b) shared priors (0.6, 0.3, 0.1). Dashed: the KL ceiling of Proposition 3.

Table A7. Post-hoc comparison of the two bets on the existing games (sequential test, $\alpha = 5\%$; for honest conditions only the share flagged applies).

Condition	Games	Joint bet: flagged	Joint bet: pair named	Likelihood ratio: flagged	Likelihood ratio: pair named
Secret plan, Haiku	40	85%	82.5%	70%	62.5%
Secret plan, Haiku (replication)	20	90%	90%	70%	70%
Secret plan, DeepSeek	40	0%	0%	32.5%	2.5%
Secret plan, DeepSeek, warned	30	0%	0%	26.7%	13.3%
Pooled-data vendor	40	50%	45%	100%	92.5%
Bid-rotation cartel	40	5%	5%	60%	40%
Secret plan, GPT-6 Luna	20	20%	20%	35%	30%
Secret plan, Gemma 3 12B	20	15%	10%	30%	10%
Secret plan, gpt-oss-20b	20	0%	0%	20%	5%
Honest, 4 × DeepSeek	150	0%	–	4.7%	–
Honest, 4 × Haiku	150	0%	–	2%	–
Honest, mixed models	150	0.7%	–	5.3%	–
Honest, Haiku (replication)	40	0%	–	0%	–
Honest, auction	100	0%	–	0%	–
Public-data vendor	60	0%	–	0%	–

Appendix D. Reproducibility and Deviations

Replays and seeds. Every reported agent-round has exactly two replays (the analysis would use the smallest count available in a game). Each condition uses its own block of episode seeds, and the main samples, pilots and the replication use disjoint blocks. Each model call's sampling seed is derived from the call type, the episode seed, the agent and the round, so a logged move and its replays never share a seed; the null resampling of each episode is seeded by its episode seed. Pilot games were excluded from every analysis.

Deviations from the plan.

- Connection errors in the first launch were handled by re-running the same seeds.
- A post-hoc statistic comparing each agent's logged move with its own replays did not improve detection and is not used.
- The stronger judge, unnamed in the plan, was run on 20 games per condition for budget reasons, with replies of up to 2,000 tokens.
- A review found that the sequential test named the pair by its final evidence instead of at the alarm round; we fixed this and recomputed all sequential results.
- The judge is not deterministic at temperature 0. On games without private messages its two views send the same prompt, and an early uncached pass drew different replies for them (AUROCs of 0.84 for Haiku and 0.47 for DeepSeek in that pass); the reported numbers use the single cached reply, which reproduces exactly.
- Table A4, the AUROC intervals and Fisher's test were added after the reviews, from the frozen statistics.
- Models that failed the pilot's format or availability checks were dropped before the main samples.

Appendix E. LLM Usage Statement

We used Claude (Anthropic) as a coding and analysis assistant: it implemented the experiment code under our direction, ran simulations, summarised related work and talk transcripts, drafted sections of this report and generated the figures from the result files. Codex (OpenAI) and separate Claude instances served as adversarial reviewers of the plan and of the report. Every number in the report is produced by scripts from the result files. We checked claims and references against their sources and take responsibility for the content.